

ECE 6524 – Deep Learning

Homework Assignment #3: Logistic Regression

(1) Logistic regression for two-class classification:

- N pairs of data points: (\mathbf{x}_i, C_i) , $i = 1, \dots, N$, where \mathbf{x}_i is the feature vector, and C_i is the binary class label (i.e., C_i is either 0 or 1;).
- Logistic regression is a (special) linear classifier: $y_i = \boldsymbol{\beta}_1^T \mathbf{x}_i + \beta_0$, but y_i is interpreted as log-odds, i.e., $y_i = \log(p / (1 - p))$, where the probability of class 1 ($C_i = 1$) is p , and the probability of class 0 ($C_i = 0$) is $1 - p$. The prediction is based on $\text{sigmoid}(y_i) = 1/[1+\exp(-y_i)]$; if $\text{sigmoid}(y_i) \geq 0.5$, \mathbf{x}_i is predicted as a data sample from class 1; otherwise, class 0.
- The likelihood function is: $\prod_{i=1, \dots, N} (p(\mathbf{x}_i))^{C_i} (1 - p(\mathbf{x}_i))^{(1 - C_i)}$, where $p(\mathbf{x}_i)$ is the probability of \mathbf{x}_i from class 1.

Prove that the log-likelihood function can be rewritten/formulated as

$\sum_{i=1, \dots, N} \{C_i (\boldsymbol{\beta}_1^T \mathbf{x}_i + \beta_0) - \log[1 + \exp(\boldsymbol{\beta}_1^T \mathbf{x}_i + \beta_0)]\}$, and **compute the gradient** of the log-likelihood function.

(2) As instructed in Homework Assignment #0, we can build a simple Python program to construct a logistic regression classifier to test on simulate data.

- Generating simulated data (two-dimensional data of two classes with some overlap; as illustrated in Fig. 1) by sampling from multivariate normal distributions.




Fig. 1. An example of simulated data: for example, data points in brown are in class 1, and data points in blue are in class 0.

- b. Build a logistic regression model using the scikit-learn library. Below is an example for building a logistic regression model:

```
from sklearn.linear_model import LogisticRegression

clf = LogisticRegression(fit_intercept=True, C = 1e15)
clf.fit(simulated_features, simulated_labels)

print clf.intercept_, clf.coef_
```

- (3) Implement a two-class logistic regression classifier with the **stochastic gradient descent** (SGD) algorithm, where the **loss function** is the log-likelihood function. Specifically, you are asked to 'replace' the `clf.fit()` method with your own implementation of the method such as the following:

```
def logistic_regression(features, class_label, num_steps, learning_rate)

    # TO DO: YOUR IMPLEMENTATION

    return weights
```

Verify that the weights estimated from SGD are in consistent with the weights obtained from the `clf.fit()` method.

- (4) Investigate the impact of learning rate on the convergence of the algorithm. (Hint: you may plot the log-likelihood over time to show the convergence.)
- (5) Report the prediction performance on the simulated data.
- (6) Apply your logistic regression classifier to some real data of your choice (like the Iris data and MNIST data (<http://yann.lecun.com/exdb/mnist/>)) in predicting any two classes. [Note that you need to extend your implementation to handle n-dimensional data.] Experiment with different choices of two classes as many as possible if you use a multiclass data set, and report the overall performance of your logistic regression classifier.
- (7) Discuss about how you can extend your implementation to classify multiclass data (i.e., the number of class > 2).

You need to prepare a **written report** (in the pdf format) including the following sections: (1) Introduction or Problem Statement, (2) Method(s) or Approach(es), (3) Experimental Results, and (4) Discussion & Conclusion.

In addition, you need to attach your **implementation codes** as separate files to the report.